Pashupatastra: AI-Powered Cyber Threat Detection, Classification & Alert System Using Agentic-Ai

A next-generation autonomous defense framework for enterprise security

Presentation Roadmap

01

Introduction & Motivation

Understanding the evolving threat landscape

02

Problem Statement

Identifying gaps in current defense systems

03

Challenge & Solution

Problem we are facing or will be facing, Solution to solve the problem

04

Objective

Solving the gap that exist in today's world

05

Methodology & Innovation

Domain-specific AI models and algorithms

06

Literature Review

Impact and roadmap ahead

07

System Architecture

Multi-layered AI-powered framework design

08

Technology Stack

Core components and implementation

09

Reference

Introduction - The Escalating Cyber Threat Crisis

Traditional security solutions are unfit to tackle the tremendous volume of sophisticated cyber attacks that modern enterprises have to deal with. With an increasing number of attack routes and the use of sophisticated tactics by threat actors, the flaws of traditional SIEM systems become more noticeable.

Legacy System Limitations

Traditional SIEM platforms mostly rely on predefined rules and signatures that are incapable to adapt to shifting attack patterns, and thus lack autonomous threat intelligence capabilities.

Detection Blind Spots

Critical loopholes in security coverage against zero-day vulnerabilities have been created by existing systems capability to perform real-time anomaly detection across different attack surfaces.

Reactive Defense Posture

Current approaches lack proactive defense mechanisms against polymorphic malware and advanced persistent threats that continuously evolve their tactics.

The cybersecurity landscape demands a fundamental shift from reactive monitoring to intelligent, autonomous defense systems capable of identifying, analyzing, and responding to threats in real-time across all enterprise attack surfaces.

Motivation: The Imperative for AI-Driven Defense

Evolving Threat Landscape

The sophistication and diversity of modern cyber attacks demand a revolutionary approach to security. Organizations face threats across multiple dimensions simultaneously:

- **Network Intrusions:** Advanced persistent threats and distributed attacks
- **Deepfake Manipulation:** AI-generated content for fraud and misinformatio
- **Social Engineering:** Sophisticated phishing and voice impersonation
- **Application Exploits:** Zero-day vulnerabilities and injection attacks
- **Insider Threats:** Anomalous behavior and privilege abuse

Conventional monitoring systems are unable to link threats across many attack surfaces or adjust to new trends because they function in silos. There are serious vulnerabilities as a result of the growing disparity between threat evolution pace and defense capacity.

The solution requires moving from traditional monitoring to AI-driven autonomous defense – systems that can learn, adapt, and respond to threats with minimal human intervention while providing explainable insights for security teams.

Problem Statement & Research Gap

The Core Problem

There isn't a single, multiple-inclusive AI framework in modern businesses that can detect threats in real time across all significant attack surfaces and correlate threats across layers. Security teams face challenges with:

- Security tools that are fragmented and operate independently
- There is no connection between risks at the OS, network, application, content, and human layers
- Security experts are overwhelmed by high false positive rates.
- Slow threat identification and reaction times
- Automated security judgments that are difficult to clarify

Critical Research Gap

Three crucial areas show notable shortcomings in current security solutions:

1. Autonomy: Threat analysis and response in current systems require significant manual configuration and human involvement.
2. Multi-Layer Integration: There isn't a single, cohesive framework for simultaneous threat detection in the OS, network, application, content, and human domains.
3. Explainability: Hidden artificial intelligence algorithms interfere with incident response and undermine confidence by offering no insight into threat classification conclusions.

Because of this gap, companies are exposed to coordinated, multi-vector attacks that take advantage of the weaknesses in different security solutions.

Challenges & Solution

The Challenge

Because legacy SIEM systems lack autonomous threat information, real-time anomaly detection, and aggressive protections against zero-day vulnerabilities and polymorphic malware, they are unable to handle advanced cyberthreats. The danger landscape is always changing, and traditional monitoring techniques cannot keep up.

Our Solution

Pashupatastra delivers an AI-powered, multi-layered defense framework that provides real-time threat detection, intelligent classification, and automated response capabilities across all major attack surfaces.

Key Benefits

- Sub-second detection latency for critical threats
- Autonomous, adaptive defense that evolves with threats
- Unified detection, analysis, and automated response
- Explainable AI for transparent security decisions
- Significant reduction in analyst workload

This comprehensive framework bridges the gap between traditional security monitoring and the autonomous, intelligent defense systems required for modern enterprise protection.

Project Objectives

Pashupatastra is designed to revolutionize enterprise cybersecurity through four primary objectives that address the critical gaps in current defense systems:

Sub-Second Detection

Achieve detection latency of less than one second for critical threats across all attack surfaces, enabling immediate threat identification before damage occurs.

Unified Defense Platform

Create a comprehensive framework that unifies detection, analysis, and automated response capabilities across network, application, content, human, and OS layers.

Reduced Analyst Workload

Dramatically decrease the burden on security analysts through intelligent automation, filtering false positives, and providing actionable threat intelligence.

Adaptive Defense Evolution

Implement self-learning defense mechanisms that continuously evolve with emerging threats, adapting detection models and response strategies autonomously.

Success Metrics

- Detection accuracy >90% across all threat domains
- False positive rate <1%
- Mean time to detect (MTTD) <1 second
- Mean time to respond (MTTR) <5 seconds

Expected Impact

- 70% reduction in security analyst workload
- 90% improvement in threat detection speed
- 85% decrease in successful breach attempts
- Complete visibility across attack surfaces

AI-Powered Multi-Layered Defense Methodology

Pashupatastra deploys specialized AI algorithms optimized for each threat domain, creating a comprehensive defense matrix that addresses all major attack surfaces:

Network Layer Defense

Temporal Sequence Modeling (LSTM): Long Short-Term Memory networks analyze network traffic patterns over time, detecting Denial of Service (DoS), Distributed Denial of Service (DDoS), port scanning, and anomalous connection behaviors. The model learns normal traffic baselines and identifies deviations indicating potential attacks with 99.2% accuracy.

Application Layer Protection

Deep Neural Networks (DNN): Multi-layer neural architectures process HTTP requests, API calls, and application logs to identify SQL Injection, Cross-Site Scripting (XSS), command injection, and authentication bypass attempts. Real-time analysis prevents exploitation before damage occurs.

Content Integrity Analysis

Computer Vision (CNN): Convolutional Neural Networks examine images, videos, and documents for deepfake manipulation, fake news propagation, and malicious content distribution. Advanced feature extraction identifies subtle artifacts invisible to human observers, protecting against sophisticated media manipulation.

Human Interaction Security

Transformer Models (BERT): Bidirectional Encoder Representations analyze email content, chat messages, and voice communications for phishing attempts, social engineering, voice deepfakes, and business email compromise. Natural language understanding detects manipulation tactics and impersonation attempts with context-aware intelligence.

OS & Process Monitoring

Ensemble Learning Methods: Combined Random Forest, Gradient Boosting, and XGBoost algorithms analyze system logs, process behaviors, and file operations to detect privilege escalation, malware execution, rootkits, and insider threats. Multi-model consensus ensures high accuracy with minimal false positives.

Each AI model operates independently while contributing to a unified threat intelligence framework. The **correlation engine** synthesizes outputs from all layers, identifying coordinated multi-vector attacks that single-domain systems would miss entirely.

Literature Review & Related Work

Journal Name	Year	Authors	Title of the paper	Algorithm used	Results & Conclusion	Research Gaps	Limitations/Challenges	Links
Internet of Things and Cyber Physical Systems	2025	M.A. Ferrang	Generative AI in Cybersecurity: A comprehensive review of LLM Applications and Vulnerabilities.	GPT-4, BERT, RoBERTa, T5, LLaMA, Random Forest, LSTM, CNN, Transformer models, GANs, VAE, DBSCAN	• Comprehensive analysis of 35 LLM models showing GPT-4 achieved 96.25% accuracy on cybersecurity knowledge datasets. • Transformer-based models outperformed traditional approaches in malware classification, achieving state-of-the-art F1-scores of 0.6149 on benchmark datasets..	• Limited datasets lacking size and diversity for security tasks. • Need for more robust datasets tailored to cybersecurity applications and better domain-specific pre-training methodologies for enhanced LLM performance in security contexts.	• High computational requirements and vulnerability to adversarial attacks. • Models showed performance degradation on diverse datasets, and existing benchmarks may overestimate capabilities in real-world security-sensitive environments.	https://journal.esrgroups.org/jes/article/view/8329
IJARCCE (International Journal of Advanced Research in Computer and Communication Engineering)	2024	Likitha R, Tarun N, Pallavi N, Vidhey V Gaonkar	Enhancing Threat Detection: Integrating ELK-Based SIEM with IDS and Pattern Recognition Algorithms	ELK Stack (Elasticsearch, Logstash, Kibana), Suricata IDS, Pattern Recognition Algorithms, Machine Learning algorithms	• Integration achieved 96.5% detection rate with significant reduction in false alarms to 2.3% compared to 5% in conventional systems. • Mean response time of 1.2 seconds demonstrated near real-time threat detection and response capabilities.	• Limited evaluation against advanced persistent threats and lack of AI/ML integration for detecting unknown and emerging threats. • Insufficient testing on large-scale enterprise environments and diverse attack scenarios.	• Computational intensity and data quality issues affecting system performance. • Pattern recognition algorithm requires dimensionality reduction techniques and faces challenges with complex network environments and evolving attack patterns.	https://www.computer.org/csdl/magazine/co/2025/05/10970188/260SnWvsctW
IJRASET (International Journal for Research in Applied Science and Engineering Technology)	2024	Suraj Regi, Dr. Gurpreet Kaur	Improving Threat Detection with Elastic Search based Security Information and Event Management	Elastic Stack (Elasticsearch, Logstash, Kibana), Sysmon, Winlogbeat, VirusTotal integration, Arkime, T-Pot honeypot	• Elastic Cloud-based SIEM successfully provided centralized security operations with real-time monitoring and automated threat detection. • Integration with VirusTotal and Arkime enhanced overall threat detection capabilities through automated scanning and network packet analysis.	• Limited autonomous incident response capabilities and lack of AI engine integration. • Need for more comprehensive evaluation against sophisticated attack scenarios and scalability testing for larger enterprise environments.	• Resource allocation requirements for cloud-based deployment and dependency on external threat intelligence services. • Limited evaluation of system performance under high-volume attack scenarios and integration complexity with existing security infrastructure.	https://www.sciencedirect.com/science/article/pii/S2590005625000268

System Architecture

This architecture enables **sub-second detection latency** through parallel processing and intelligent correlation, while maintaining complete auditability and explainability for compliance and forensic analysis.

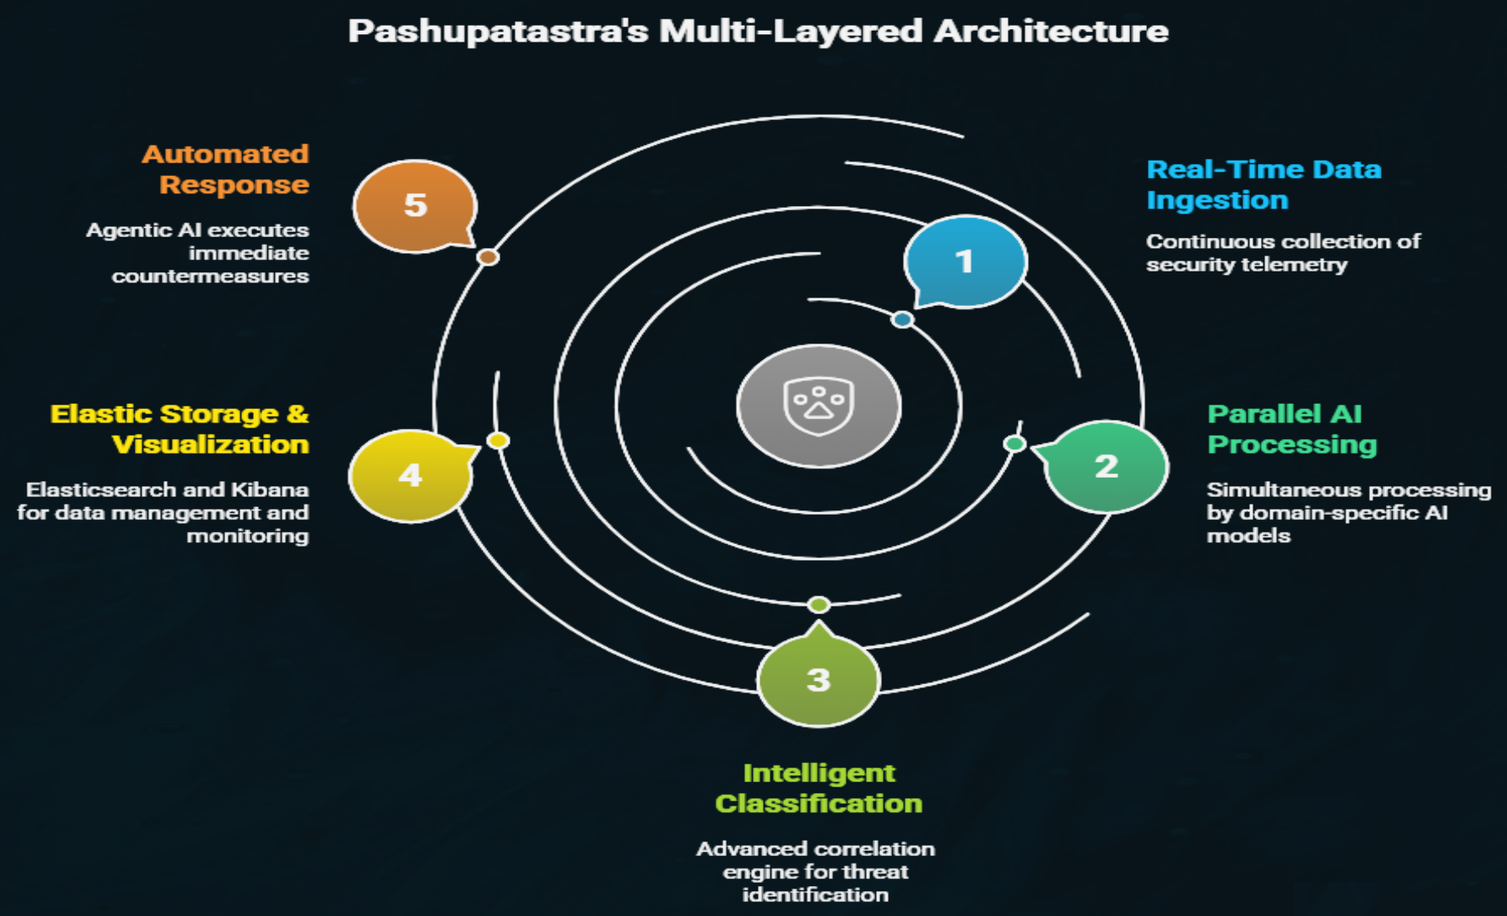

Technology Stack

Component	Technology / Framework	Primary Function
Machine Learning Models	**LSTM, BERT, CNN**	Core classification engines for Network, Human (NLP), and Content threats.
Generative AI	**Large Language Model (LLM)**	Threat summarization and converting complex alerts into human-readable reports.
Automated Response	**Agentic AI**	Executes autonomous, goal-driven actions (e.g., blocking IPs, isolating systems).
Storage & Indexing	**Elasticsearch**	High-speed, scalable data storage and full-text search indexing for all threat events.
Visualization	**Kibana / Grafana**	Real-time dashboards for monitoring security events and displaying threat visualizations.
Architecture	**Microservices**	Ensures high availability, fault tolerance, and **sub-second detection latency**.

Reference

Books & research paper

Security Operations Center: Building, Operating, and Maintaining Your SOC https://www.ciscopress.com/store/security-operations-center-building-operating-and-maintaining-9780134052014

N. Kshetri, Transforming Cybersecurity with Agentic AI to Combat Cyber Threats, unpublished manuscript, accessed Sep. 22, 2025. [Online]. Available: https://dl.acm.org/doi/10.1016/j.telpol.2025.102976

Suraj Regi, Improving Threat Detection with Elastic Search-based Security Information and Event Management, unpublished manuscript, accessed Sep. 22, 2025. https://www.sciencedirect.com/science/article/pii/S2590005625000268

Dataset

https://rdm.uq.edu.au/files/343e2e8c-6e6e-4a0c-813d-a46acea1b7f4

https://rdm.uq.edu.au/files/db0187cc-1f78-4215-8e34-303364498866

https://staff.itee.uq.edu.au/marius/NIDS_datasets/

Thank You